

Fred Smith on AI

A Panel of Experts Discussion



RICHARD RUMELT

AUG 18, 2026



I have been writing the biographies of an expert panel. Each expert brings their own peculiar set of skills, habits of mind, and life experience—like a council of elders gathered around a fire, each with a different story to tell. The panel I am writing about (book to come) is not available for hire; most of its members are long gone. Yet if you read closely enough, you can absorb their ways of thinking, their rules of thumb, and try on their perspectives as if donning a new set of glasses.

In doing so, you can step outside the narrow perimeter of your own habits and see the problem afresh.

My wife, Kate, provided the original stimulus for the panel concept. Teaching business strategy by the case method at the University of Michigan, she developed a methodology she named “create-destroy.” Facing a hard problem, she explained, is uncomfortable, even stressful. When we see a potential solution, we grasp it tightly; the relief of having something to hold onto is palpable. There may be other, better solutions to our problem, but to find them we would have to let go of the one in our hands. Not just *look* elsewhere: the gravitational pull of a first-born idea is too strong for that. We have to actively reject or *destroy* the original idea to make room for a better one.

Practicing Kate’s discipline, I found myself turning more and more to a particular kind of outside help. I summon a *virtual* panel of experts—people whose judgment I trust—and carry them with me, like a pocketful of lucky stones. My internal dialogue with them is a running conversation, sometimes a debate, that pokes holes in my own ideas and sparks new ones.

As I wrote about Fred Smith, the entrepreneur behind Federal Express, I found myself impressed all over again by his achievement. He changed the very anatomy of the air freight business. As Rockefeller did with petroleum and Swift with meatpacking, he changed not just the product but the system: the architecture of an industry.

As I finished my note on Smith, I found myself wondering what he would make of AI. Is this new thing being delivered over a system built for a different era, a different kind of good entirely? My old habits as an electrical engineer took over, and I started to pull at the threads.

Today, AI is being delivered over the web, riding on HTTP requests and responses, with the browser as its window to the world and cloud data centers as its engine room. The protocol is REST and JSON, the user

experience is the browser window, and the model must reach through SaaS silos just to touch the data layer. None of these pieces were built for this purpose. HTTP was meant for fetching documents, the browser for reading them, REST for simple operations on structured records. The chat window itself began as a humble customer-support ticket form. The SaaS API was designed to repeat the same steps, not to let a model reason about the user's world. Even the GPU, the one part that fits the new workload like a glove, started life rasterizing triangles for video games. Everything else is a retrofit, a patchwork of yesterday's tools pressed into service for today's ambitions.

What does the AI user really want, from end to end? A partner that sticks around. One that remembers, that holds context across sessions and applications, that acts on your behalf with your own authority, that delivers outcomes instead of just spitting out tokens, and that lives with you rather than waiting to be summoned. But where does this chain break? At every point where the model must squeeze through an interface built for a human: the browser tab that vanishes, the OAuth flow, the form field, the file dialog, the ticket queue, the RAG pipeline that each morning re-embeds the memory the model should have kept all along. Each is a link in a chain forged before AI was even imagined.

A purpose-built AI transport would not be a chat window tacked onto the side of an application. It would be a runtime layer running underneath, woven into the fabric of the system itself. Persistent memory would not be a hack or a workaround; it would be as fundamental as breathing.

MCP is the first serious attempt at a new transport, but it still rides on JSON-RPC over HTTP. In that sense, it is more like Smith's first Falcon 20—a clever aircraft that slipped through a regulatory loophole—than the plane the system truly needs. The loophole, though, is different this time. Smith's Falcon 20 exemption was a regulatory gift, his alone for a while. MCP, by contrast, is an open standard; anyone can fly it as soon as it lands. The real gate now is not regulation but the incumbents'

terms of service, API pricing, and platform policies—the modern echo of the CAB’s route authority. If we priced by value delivered, not by tokens counted, and built a trust layer for real actions on the user’s behalf, not just retrofitted web authentication, we might finally have the right system.

Perhaps the winner will be whoever owns the chain from end to end, just as Smith won by refusing to ride on the back of passenger routes. Apple is making its play at the device layer with Apple Intelligence and Private Cloud Compute. Microsoft is staking its claim at the OS layer with Copilot. Anthropic is inching toward it with computer use, MCP, and the Agent SDK.

The web-native incumbents will be hard to unseat as long as the retrofit is good enough—just as passenger carriers ruled air freight for a decade after Smith’s paper. That is the era we are living through now: the passenger-airline phase of AI, a patchwork solution that everyone accepts simply because the real system has not yet been revealed.

But then my strategist self started to argue with the imagined ghost of Smith. He could own the chain end to end because his business was about moving a package—the customer stayed outside the box. AI is different. Here, the customer is inside the transaction from start to finish. What is being transported is not a package, but a relationship with the interior of a person’s life: their worries, their money, their hopes. The question shifts from chain of custody to chain of trust, and that chain is longer, more fragile, and far harder to forge.

The professions that have historically been allowed into that interior—the family doctor, the trusted lawyer, the private banker, the priest, the longtime accountant, the family office—each earned trust through a mix of personal continuity, enforceable privilege, and professional discipline. That trust let the practitioner fade into the background, their identity less important than their role. They had clear boundaries and limits on what could be done with what they knew. None of this exists for AI

providers today. Your conversation can be trained on, stored, subpoenaed, or mined for product tweaks. There is no fiduciary duty, no malpractice bond, no professional body to strike off a bad actor. Trust is being asked for without any of the scaffolding that once supported it, and the asking is running far ahead of the earning.

The people who solve this trust problem probably won't be the model builders. It may be the old trust institutions—the professional firm, the private bank, the family doctor, the family office—using AI as a tool within the fiduciary posture they already know. Or it may be new institutions, built from the ground up around AI-specific fiduciary discipline, licensed and bonded as medicine is. Apple's device-local architecture is one wager: dissolve the trust problem by cutting out the third party. The professional-fiduciary model is another: keep the third party, but make it accountable, as lawyers and doctors are.

Maybe the real winner will be whoever earns that trust, forging the deeper chain. Owning the transport is just the ante—the price of admission to the real game.